

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

```
In [2]: data = pd.read_csv("C:/Users/anjali/Desktop/diabetes.csv", sep=",")
```

```
In [3]: data.shape
```

```
Out[3]: (768, 9)
```

```
In [4]: data.head(5)
```

```
Out[4]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFui
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	

```
In [5]: data.columns ## to see all the columns present in the data
```

```
Out[5]: Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness',
              'Insulin',
              'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],
              dtype='object')
```

```
In [6]: data.dtypes ##to see the datatypes of each col
```

```
Out[6]: Pregnancies          int64
Glucose          int64
BloodPressure    int64
SkinThickness    int64
Insulin          int64
BMI              float64
DiabetesPedigreeFunction  float64
Age              int64
Outcome          int64
dtype: object
```

In [7]: data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Pregnancies                          768 non-null    int64
1   Glucose                              768 non-null    int64
2   BloodPressure                        768 non-null    int64
3   SkinThickness                       768 non-null    int64
4   Insulin                              768 non-null    int64
5   BMI                                  768 non-null    float64
6   DiabetesPedigreeFunction             768 non-null    float64
7   Age                                  768 non-null    int64
8   Outcome                              768 non-null    int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

In [8]: data.describe() *##using numpy calculate all math functions*

Out[8]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	Di
count	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	
mean	3.845052	120.894531	69.105469	20.536458	79.799479	31.992578	
std	3.369578	31.972618	19.355807	15.952218	115.244002	7.884160	
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	1.000000	99.000000	62.000000	0.000000	0.000000	27.300000	
50%	3.000000	117.000000	72.000000	23.000000	30.500000	32.000000	
75%	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000	
max	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	

In [9]: data.isna().sum() *##to check null values and their sums*

Out[9]:

Pregnancies	0
Glucose	0
BloodPressure	0
SkinThickness	0
Insulin	0
BMI	0
DiabetesPedigreeFunction	0
Age	0
Outcome	0
dtype:	int64

```
In [10]: data.isna().any()
```


```
Out[10]: Pregnancies      False
          Glucose          False
          BloodPressure    False
          SkinThickness     False
          Insulin           False
          BMI               False
          DiabetesPedigreeFunction  False
          Age               False
          Outcome           False
          dtype: bool
```

Data analysis and visualization


```
In [11]: data.Outcome.value_counts() ##outcome is the classification colmn and value_counts count the no of discrete outcomes.
```

```
Out[11]: 0      500
          1      268
          Name: Outcome, dtype: int64
```


```
In [12]: sns.countplot(data.Outcome, data=data)
          plt.show() ## plot the above data using seaborn
```



```
In [13]: sns.distplot(data['Age'], kde=True)
plt.show()  ##KDE find the density of the variable(age) using probab
           ility density function and gaussian algo
```




```
In [14]: sns.distplot(data['Insulin'], kde=True)
plt.show()
```

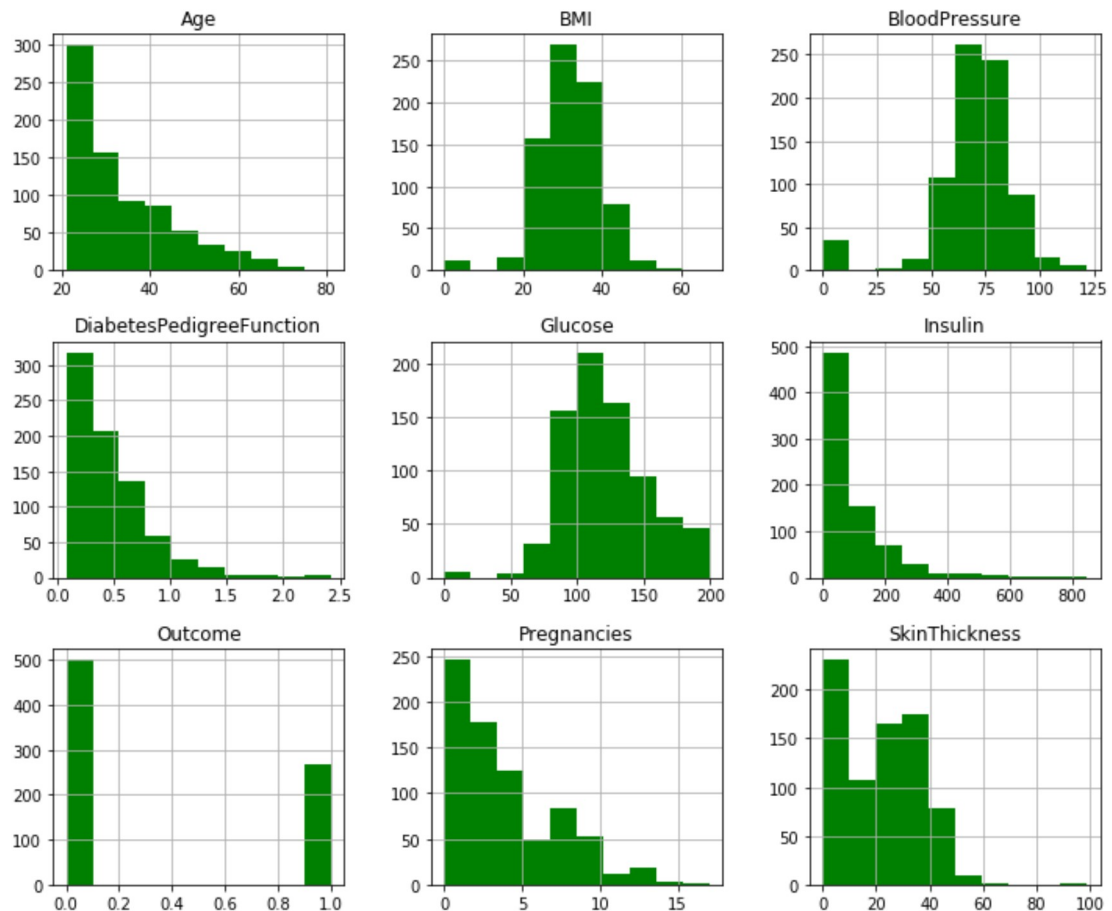


```
In [15]: correlation = data.corr()
plt.figure(figsize = (12,10))
sns.heatmap(correlation,annot = True) ##correlation used to find the dependencies of each independent variable(columns) on each other
```


```
Out[15]: <matplotlib.axes._subplots.AxesSubplot at 0x205360da548>
```



```
In [16]: data.hist(figsize = (12,10),color = 'green')
plt.show() ##hist plot for each column
```



```
In [17]: plt.hist(data.Age,bins = [10,20,30,40,50,60,70,80,90])
plt.show() ## bins is the interval between which the histogram needs to be plotted
```




```
In [18]: data.groupby("Outcome").mean() ##mean of each variable according to outcome
```


```
Out[18]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI
Outcome						
0	3.298000	109.980000	68.184000	19.664000	68.792000	30.304200
1	4.865672	141.257463	70.824627	22.164179	100.335821	35.142537


```
In [19]: # box plot between outcome and insulin
sns.boxplot(x='Outcome', y='Insulin', data=data)
plt.show()
```




```
In [20]: # scatter plot between glucose and Blood pressure with regression line
sns.regplot(x="Glucose", y="BloodPressure", data=data)
plt.show()
```



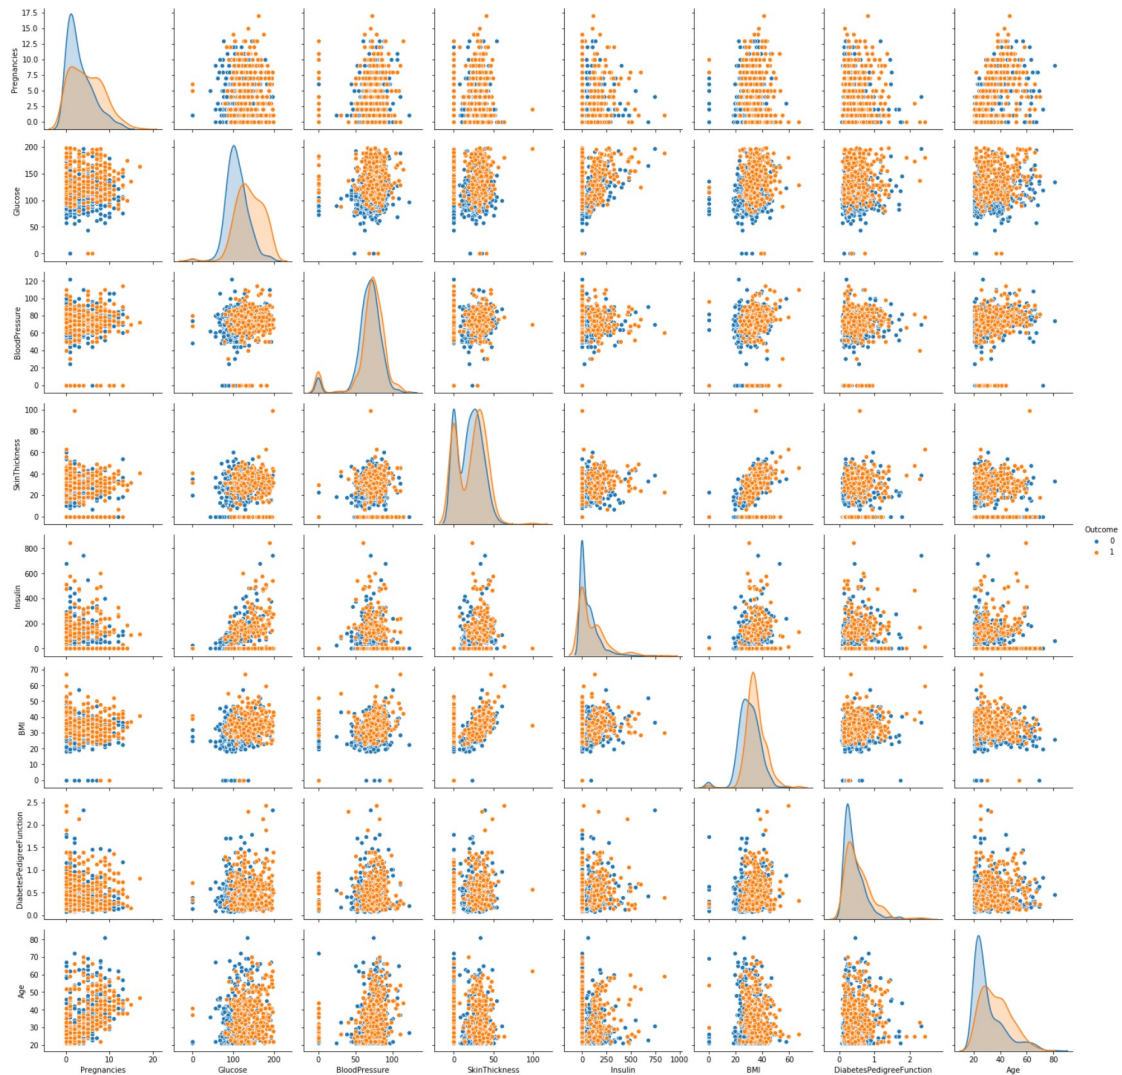
```
In [21]: sns.regplot(x='BMI', y= 'Glucose', data=data)
plt.show()
```



```
In [22]: sns.scatterplot(x='Glucose', y= 'Insulin', data=data) # scatter plot
plt.show()
```




```
In [23]: sns.pairplot(data,hue='Outcome')
plt.show()
```



Now we are done with EDA and we'll start predicting using various algorithm

```
In [24]: X = data.iloc[:, :8] ##to locate data
Y = data.iloc[:, 8]
```

```
In [25]: from sklearn.svm import SVC ##using SKlearn library to import and use various ML algorithms
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, classification_report
```

```
In [26]: scaler = StandardScaler()
standardized_data = scaler.fit_transform(X) ##scaling the data using PCA
```

```
In [27]: pd.DataFrame(standardized_data,columns=['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
        'BMI', 'DiabetesPedigreeFunction', 'Age']).head() ##Dimension reduction is performed and dataset is scaled
```

Out[27]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPe
0	0.639947	0.848324	0.149641	0.907270	-0.692891	0.204013	
1	-0.844885	-1.123396	-0.160546	0.530902	-0.692891	-0.684422	
2	1.233880	1.943724	-0.263941	-1.288212	-0.692891	-1.103255	
3	-0.844885	-0.998208	-0.160546	0.154533	0.123302	-0.494043	
4	-1.141852	0.504055	-1.504687	0.907270	0.765836	1.409746	

```
In [28]: train_x,test_x,train_y,test_y = train_test_split(standardized_data,
        Y,test_size = 0.3,random_state = 32) ##splitting a dataset
```

```
In [29]: print("Shape of train x: ",train_x.shape)
        print("Shape of train y: ",train_y.shape)

        print("Shape of test x: ",test_x.shape)
        print("Shape of test y: ",test_y.shape)
```

```
Shape of train x: (537, 8)
Shape of train y: (537,)
Shape of test x: (231, 8)
Shape of test y: (231,)
```

SVM MODEL

```
In [30]: # defining our model
        model = SVC()
        model.fit(train_x,train_y)
```

```
Out[30]: SVC(C=1.0, break_ties=False, cache_size=200, class_weight=None, coef0=0.0,
        decision_function_shape='ovr', degree=3, gamma='scale', kernel='rbf',
        max_iter=-1, probability=False, random_state=None, shrinking=True,
        tol=0.001, verbose=False)
```

```
In [31]: # training accuracy
        X_train_predicted = model.predict(train_x)
        print("Accuracy of training data is: ",accuracy_score(train_y,X_train_predicted))
```

```
Accuracy of training data is: 0.8342644320297952
```

```
In [32]: # prediction for test data
        predicted = model.predict(test_x)
```

```
In [33]: svm_test_score = accuracy_score(test_y,predicted)
svm_test_score
```

```
Out[33]: 0.7489177489177489
```

```
In [34]: print(classification_report(test_y,predicted))
```


	precision	recall	f1-score	support
0	0.77	0.87	0.82	147
1	0.70	0.54	0.61	84
accuracy			0.75	231
macro avg	0.73	0.70	0.71	231
weighted avg	0.74	0.75	0.74	231

```
In [35]: import seaborn as sns
cm = confusion_matrix(test_y,predicted)
ax = sns.heatmap(cm/np.sum(cm), annot=True,
                  fmt='.2%', cmap='Blues')

ax.set_title('Confusion Matrix');
ax.set_xlabel('\nPredicted Values');
ax.set_ylabel('Actual Values ');

## Ticket labels - List must be in alphabetical order
ax.xaxis.set_ticklabels(['False','True'])
ax.yaxis.set_ticklabels(['False','True'])

## Display the visualization of the Confusion Matrix.
plt.show()
```



LOGISTIC REGRESSION

```
In [36]: from sklearn.linear_model import LogisticRegression
lr = LogisticRegression()
```

```
In [37]: lr.fit(train_x,train_y)
```

```
Out[37]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                             intercept_scaling=1, l1_ratio=None, max_iter=100,
                             multi_class='auto', n_jobs=None, penalty='l2',
                             random_state=None, solver='lbfgs', tol=0.0001,
                             verbose=0,
                             warm_start=False)
```

```
In [38]: # training accuracy
X_train_predicted = lr.predict(train_x)
print("Accuracy of training data is: ",accuracy_score(train_y,X_train_predicted))
```

```
Accuracy of training data is: 0.7802607076350093
```

```
In [39]: # prediction for test data
predicted = lr.predict(test_x)
lr_test_score = accuracy_score(test_y,predicted)
lr_test_score
```

```
Out[39]: 0.7662337662337663
```

```
In [40]: print(classification_report(test_y,predicted))
```


	precision	recall	f1-score	support
0	0.78	0.88	0.83	147
1	0.73	0.56	0.64	84
accuracy			0.77	231
macro avg	0.76	0.72	0.73	231
weighted avg	0.76	0.77	0.76	231

```
In [41]: import seaborn as sns
cm = confusion_matrix(test_y, predicted)
ax = sns.heatmap(cm/np.sum(cm), annot=True,
                 fmt='.2%', cmap='Blues')

ax.set_title('Confusion Matrix');
ax.set_xlabel('\nPredicted Values')
ax.set_ylabel('Actual Values ');

## Ticket labels - List must be in alphabetical order
ax.xaxis.set_ticklabels(['False', 'True'])
ax.yaxis.set_ticklabels(['False', 'True'])

## Display the visualization of the Confusion Matrix.
plt.show()
```



RANDOM FOREST

```
In [42]: from sklearn.ensemble import RandomForestClassifier
```

```
In [43]: rf = RandomForestClassifier(n_estimators = 50, max_leaf_nodes = 66, max_samples = 66)
rf.fit(train_x, train_y)
```

```
Out[43]: RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
                               criterion='gini', max_depth=None, max_features='auto',
                               max_leaf_nodes=66, max_samples=66,
                               min_impurity_decrease=0.0, min_impurity_split=None,
                               min_samples_leaf=1, min_samples_split=2,
                               min_weight_fraction_leaf=0.0, n_estimators=50,
                               n_jobs=None, oob_score=False, random_state=None,
                               verbose=0, warm_start=False)
```

```
In [44]: # training accuracy
X_train_predicted = rf.predict(train_x)
print("Accuracy of training data is: ",accuracy_score(train_y,X_train_predicted))
```

Accuracy of training data is: 0.8212290502793296

```
In [45]: # prediction for test data
predicted = rf.predict(test_x)
rf_test_score = accuracy_score(test_y,predicted)
rf_test_score
```

Out[45]: 0.7619047619047619

```
In [46]: print(classification_report(test_y,predicted))
```


	precision	recall	f1-score	support
0	0.76	0.92	0.83	147
1	0.77	0.49	0.60	84
accuracy			0.76	231
macro avg	0.77	0.70	0.71	231
weighted avg	0.76	0.76	0.75	231

```
In [47]: import seaborn as sns
cm = confusion_matrix(test_y, predicted)
ax = sns.heatmap(cm/np.sum(cm), annot=True,
                 fmt='.2%', cmap='Blues')

ax.set_title('Confusion Matrix');
ax.set_xlabel('\nPredicted Values')
ax.set_ylabel('Actual Values ');

## Ticket labels - List must be in alphabetical order
ax.xaxis.set_ticklabels(['False', 'True'])
ax.yaxis.set_ticklabels(['False', 'True'])

## Display the visualization of the Confusion Matrix.
plt.show()
```



DECISION TREE

```
In [48]: from sklearn.tree import DecisionTreeClassifier
```

```
In [49]: dt = DecisionTreeClassifier(criterion='entropy', min_samples_split =
6, min_samples_leaf = 25)
dt.fit(train_x, train_y)
```

```
Out[49]: DecisionTreeClassifier(ccp_alpha=0.0, class_weight=None, criterion
='entropy',
                                max_depth=None, max_features=None, max_leaf
_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_spl
it=None,
                                min_samples_leaf=25, min_samples_split=6,
                                min_weight_fraction_leaf=0.0, presort='depr
ecated',
                                random_state=None, splitter='best')
```

```
In [50]: # training accuracy
X_train_predicted = dt.predict(train_x)
print("Accuracy of training data is: ",accuracy_score(train_y,X_train_predicted))
```

Accuracy of training data is: 0.7858472998137802

```
In [51]: # prediction for test data
predicted = dt.predict(test_x)
dt_test_score = accuracy_score(test_y,predicted)
dt_test_score
```

Out[51]: 0.7575757575757576

```
In [52]: print(classification_report(test_y,predicted))
```


	precision	recall	f1-score	support
0	0.81	0.82	0.81	147
1	0.67	0.65	0.66	84
accuracy			0.76	231
macro avg	0.74	0.74	0.74	231
weighted avg	0.76	0.76	0.76	231


```
In [53]: import seaborn as sns
cm = confusion_matrix(test_y, predicted)
ax = sns.heatmap(cm/np.sum(cm), annot=True,
                 fmt='.2%', cmap='Blues')

ax.set_title('Confusion Matrix');
ax.set_xlabel('\nPredicted Values')
ax.set_ylabel('Actual Values ');

## Ticket labels - List must be in alphabetical order
ax.xaxis.set_ticklabels(['False', 'True'])
ax.yaxis.set_ticklabels(['False', 'True'])

## Display the visualization of the Confusion Matrix.
plt.show()
```



KNN

```
In [54]: from sklearn.neighbors import KNeighborsClassifier

neighbors = [1, 2, 3, 5, 7, 9, 10, 12, 15, 19, 21]
test_scores = []
train_scores = []

for i in neighbors:
    knn = KNeighborsClassifier(i)
    knn.fit(train_x, train_y)

    train_scores.append(knn.score(train_x, train_y))
    test_scores.append(knn.score(test_x, test_y))
```

```
In [55]: ind=np.argmax(test_scores)
```

```
In [56]: # coressponding train and test scores
print("Test score: ",test_scores[ind])
print("Train score: ",train_scores[ind])
```

```
Test score:  0.7662337662337663
Train score:  0.7690875232774674
```

```
In [57]: # prediction for test data
predicted = knn.predict(test_x)
knn_test_score = accuracy_score(test_y,predicted)
knn_test_score
```

```
Out[57]: 0.7619047619047619
```

```
In [58]: print(classification_report(test_y,predicted))
```


	precision	recall	f1-score	support
0	0.76	0.90	0.83	147
1	0.75	0.51	0.61	84
accuracy			0.76	231
macro avg	0.76	0.71	0.72	231
weighted avg	0.76	0.76	0.75	231

```
In [59]: import seaborn as sns
cm = confusion_matrix(test_y, predicted)
ax = sns.heatmap(cm/np.sum(cm), annot=True,
                 fmt='.2%', cmap='Blues')

ax.set_title('Confusion Matrix');
ax.set_xlabel('\nPredicted Values')
ax.set_ylabel('Actual Values ');

## Ticket labels - List must be in alphabetical order
ax.xaxis.set_ticklabels(['False', 'True'])
ax.yaxis.set_ticklabels(['False', 'True'])

## Display the visualization of the Confusion Matrix.
plt.show()
```



MODELS COMPARING

```
In [60]: # comparing our all models
models = {"SVM":svm_test_score, "Logistic Regression":lr_test_score,
          "Random Forest":rf_test_score, "Decision Tree": dt_test_score,
          "KNN":knn_test_score}
model = pd.DataFrame({"Models":["SVM", "Logistic Regression", "Random Forest",
                                "Decision Tree", "KNN"],
                      "Score": [svm_test_score, lr_test_score, rf_test_score,
                                dt_test_score, knn_test_score]}, index = np.arange(1,6))
model.head(5)
```

Out[60]:

	Models	Score
1	SVM	0.748918
2	Logistic Regression	0.766234
3	Random Forest	0.761905
4	Decision Tree	0.757576
5	KNN	0.761905

```
In [61]: model.sort_values(by='Score', ascending = False)
        ## Random Forest is our best model
```

Out[61]:

	Models	Score
2	Logistic Regression	0.766234
3	Random Forest	0.761905
5	KNN	0.761905
4	Decision Tree	0.757576
1	SVM	0.748918

In []: